

Agentic RAG System - System Design Document

Author: Kiran Date: February 2026

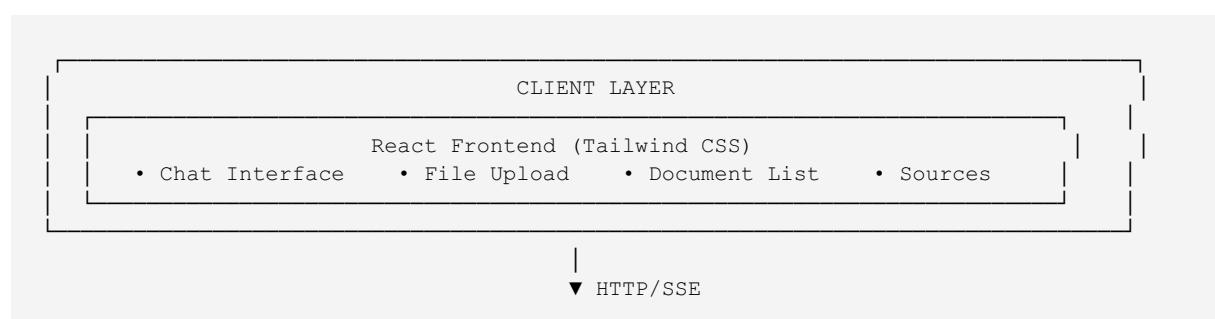
1. Executive Summary

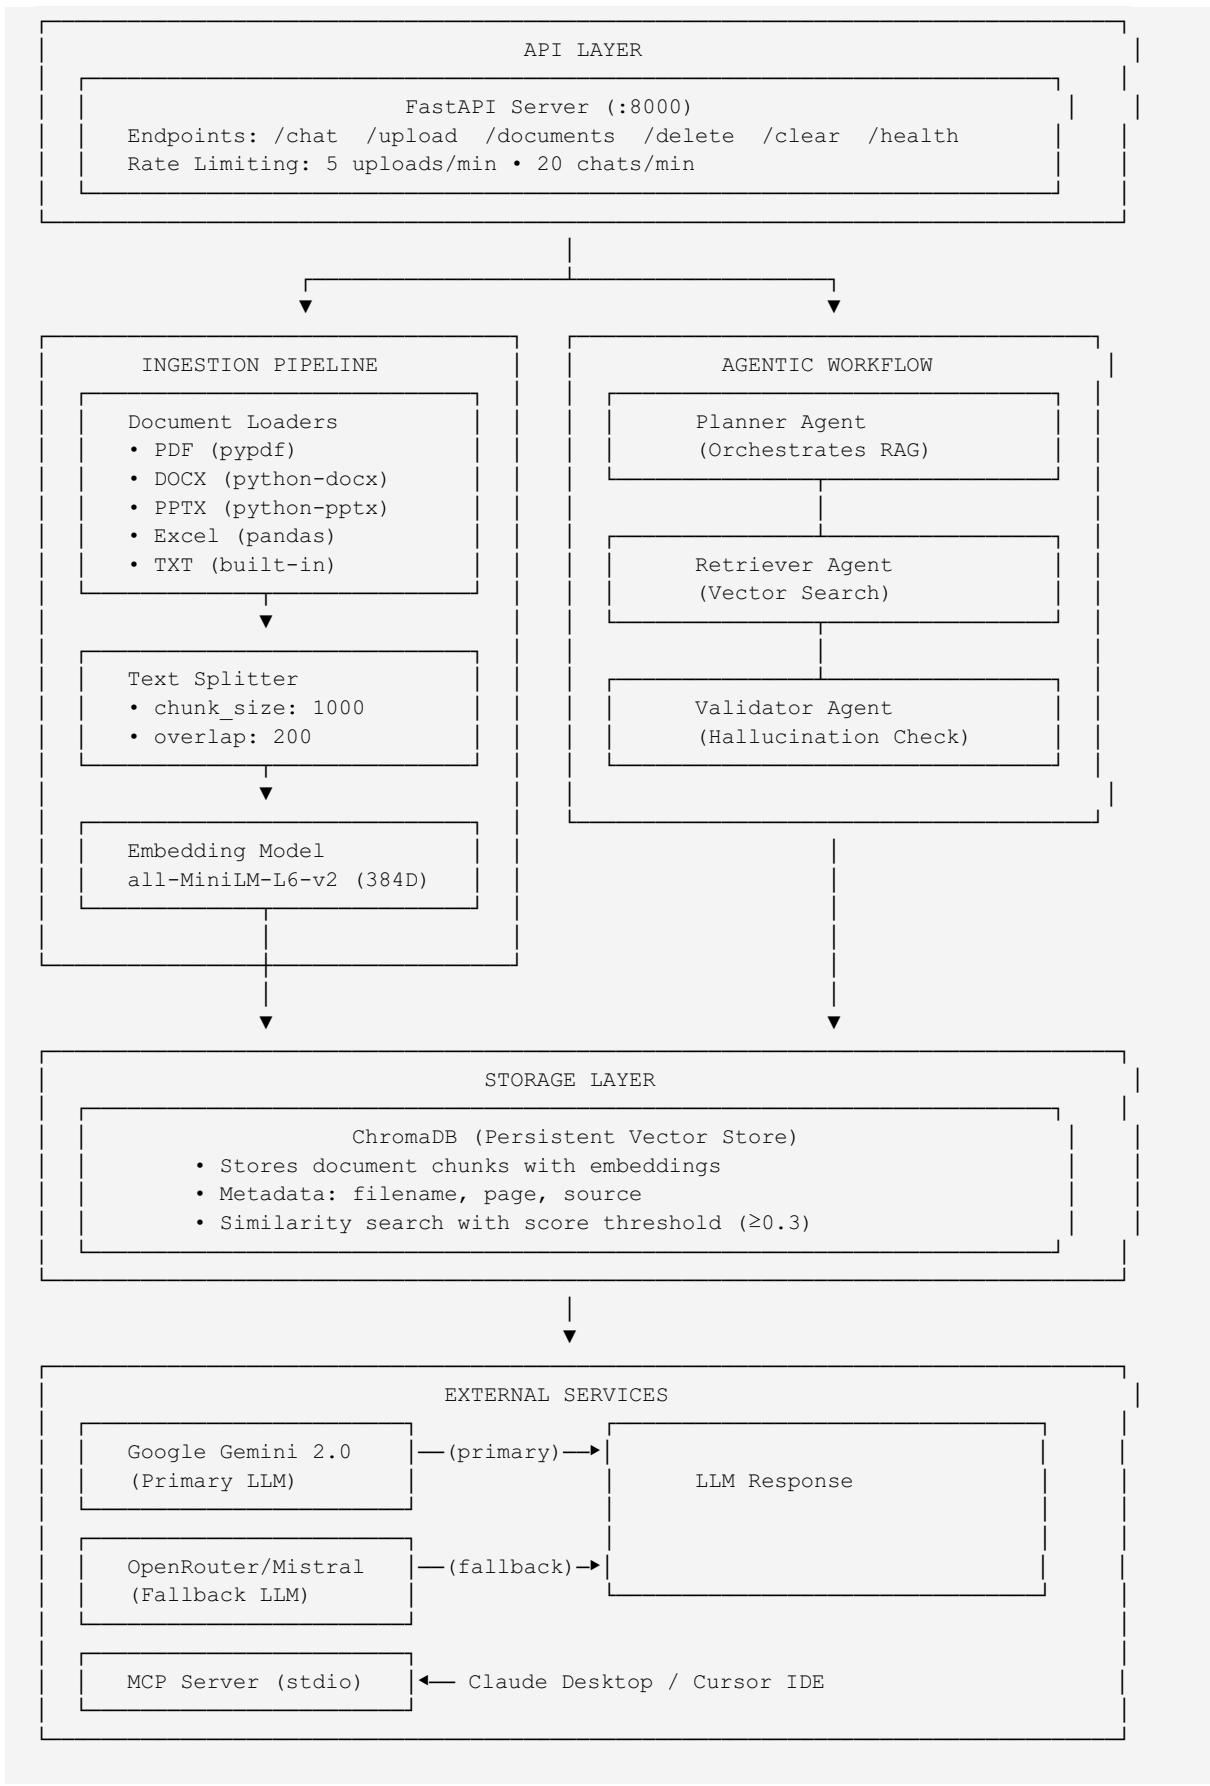
This document describes the architecture of an **Agentic RAG (Retrieval-Augmented Generation) System** that intelligently retrieves and answers questions from multi-format documents. The system combines semantic search, LLM reasoning, and agentic workflows to provide accurate, citation-backed responses.

Key Capabilities

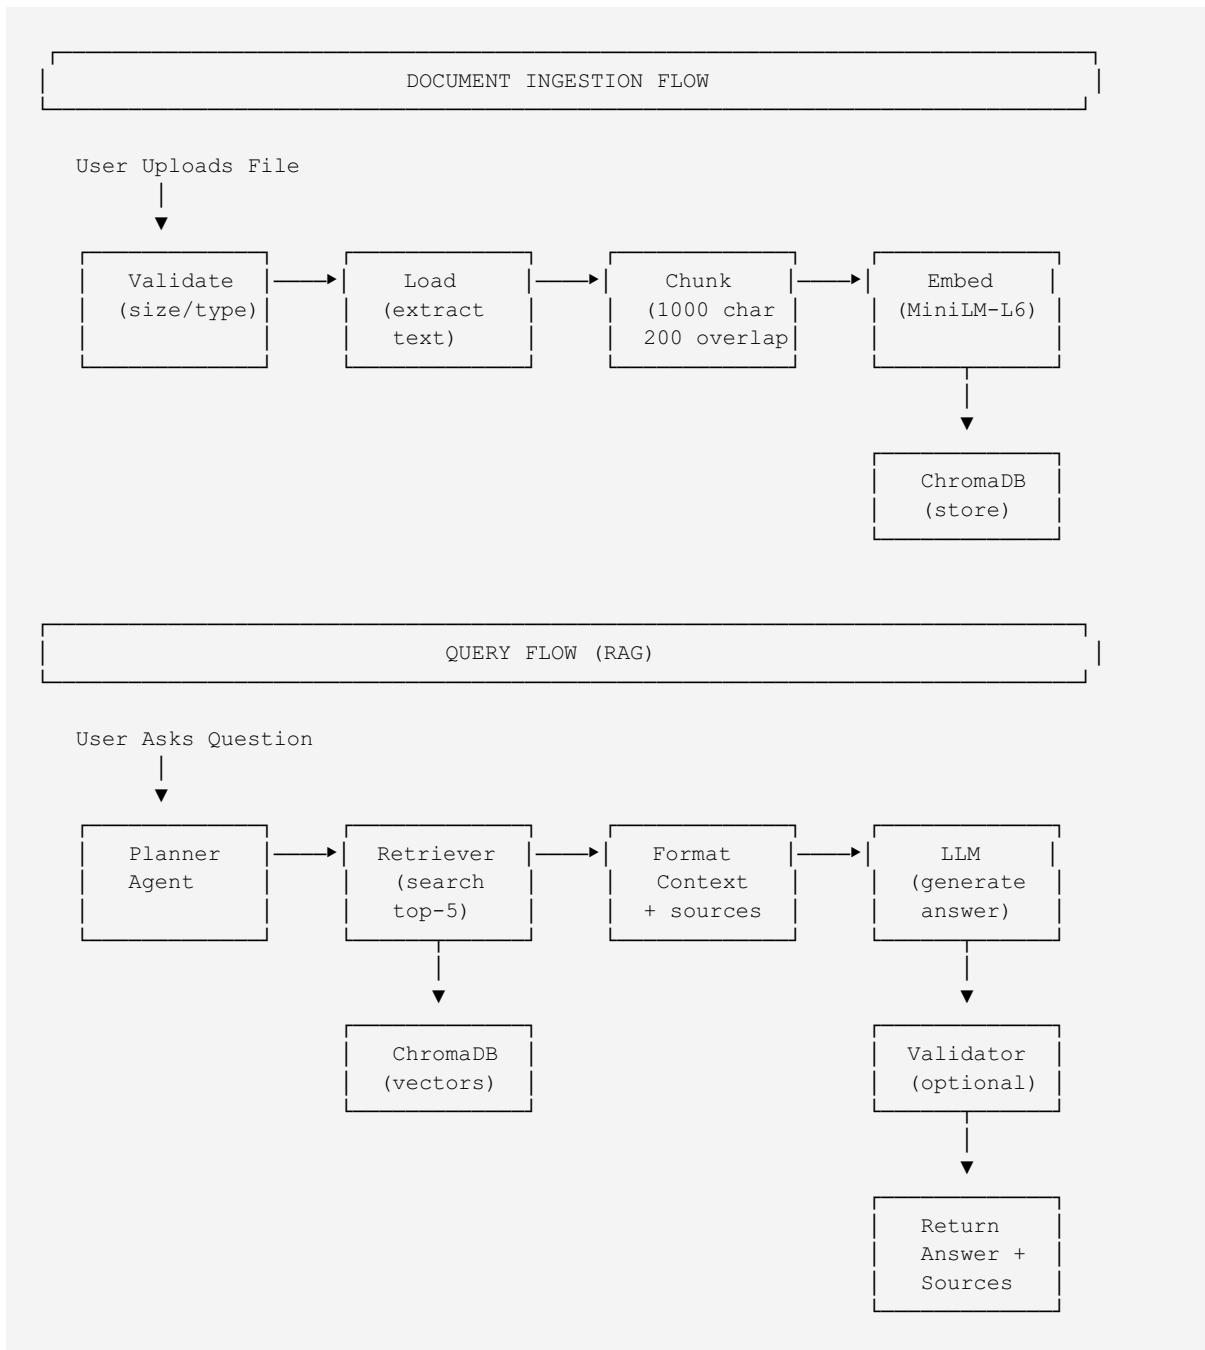
- Multi-format document ingestion (PDF, DOCX, PPTX, Excel, TXT)
 - Semantic vector search using ChromaDB
 - Agentic query processing with planning and validation
 - Real-time streaming responses via SSE
 - MCP Protocol integration for external AI tool connectivity
 - Intelligent LLM fallback strategy (Gemini → OpenRouter)
-

2. System Architecture Diagram





3. Data Flow Diagram



4. Agentic Workflow Design

The system uses a **Plan-Retrieve-Generate-Validate** pattern:

Step	Agent	Action
1	Planner	Receives user query, coordinates the workflow
2	Retriever	Searches ChromaDB for top-5 relevant chunks
3	Planner	Formats context with source citations
4	Planner	Sends context + query to LLM (Gemini/OpenRouter)
5	Validator	(Optional) Checks answer is grounded in sources
6	Planner	Returns answer + sources to frontend

LLM Fallback Strategy: - Primary: Google Gemini 2.0 Flash (best reasoning, free tier)
- Fallback: OpenRouter Mistral (if Gemini fails)

5. Context Construction Strategy

- Semantic Search:** Query embedding → ChromaDB similarity search → Top-5 chunks
 - Score Filtering:** Only chunks with similarity ≥ 0.3 are included
 - Metadata Preservation:** Each chunk includes filename, page number, source
 - Prompt Augmentation:** Context injected into system prompt with citation IDs
 - Citation Mapping:** LLM references source IDs → Frontend displays verified references
-

6. Technology Choices & Rationale

Component	Technology	Rationale
Frontend	React + Tailwind	Modern UI, responsive, good DX
Backend	FastAPI	

Component	Technology	Rationale
		Async-native, auto OpenAPI docs, Pydantic validation
Vector DB	ChromaDB	Zero-setup, persistent, ideal for < 10k docs
Embeddings	all-MiniLM-L6-v2	Local (no rate limits), 384D, fast inference
Primary LLM	Gemini 2.0 Flash	Best reasoning, large context window, free tier
Fallback LLM	Mistral (OpenRouter)	High availability, cost-effective
Streaming	Server-Sent Events	Native browser support, real-time UX

7. Key Design Decisions

- Streaming SSE Responses:** Real-time "thinking" states improve UX
- Multi-Format Loaders:** Custom loaders for Excel (row-by-row) and PPTX (slide-by-slide)
- Windows-Optimized ChromaDB:** `allow_reset=True` prevents file-lock issues
- Modular Architecture:** Loaders, embeddings, and vector DB can be swapped independently
- Rate Limiting:** Protects against API abuse (5 uploads/min, 20 chats/min)
- MCP Integration:** Exposes search tool for Claude Desktop/Cursor integration

8. Limitations

Limitation	Description	Future Solution
No Image Analysis	Only extracts text from documents	Integrate Vision LLM
Single-Tenant	All documents in one collection	Add user namespaces
Scale Ceiling	ChromaDB optimal for < 10k files	Migrate to Milvus
No OCR	Scanned PDFs not supported	Add Tesseract preprocessing
Sequential Workflow	Single plan-retrieve-generate cycle	Add iterative self-correction

9. API Reference

Endpoint	Method	Description
/chat	POST	Submit question, get answer with sources
/chat/stream	POST	Streaming SSE response
/upload	POST	Upload documents (PDF, DOCX, PPTX, XLSX, TXT)
/documents	GET	List all ingested documents
/delete/{filename}	DELETE	Remove a document

Endpoint	Method	Description
/clear	POST	Clear all data
/health	GET	Health check

This system demonstrates agentic RAG workflows with production-ready architecture for intelligent document Q&A.